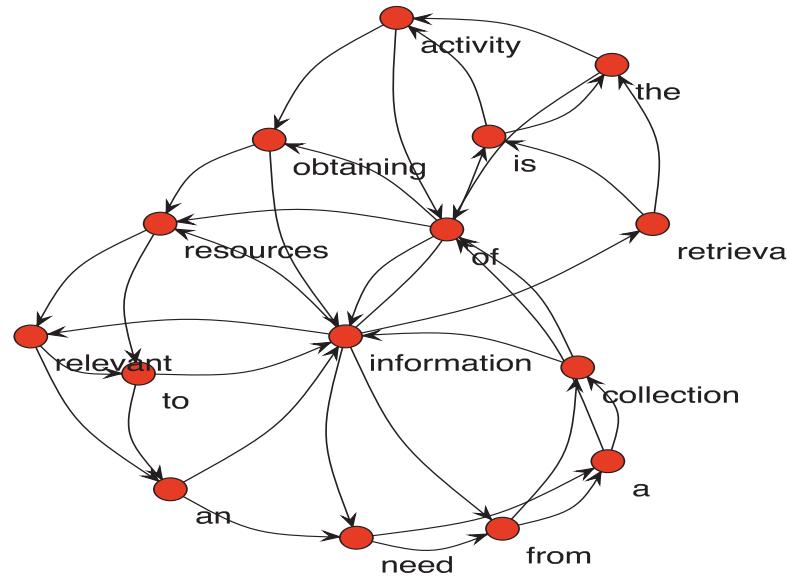


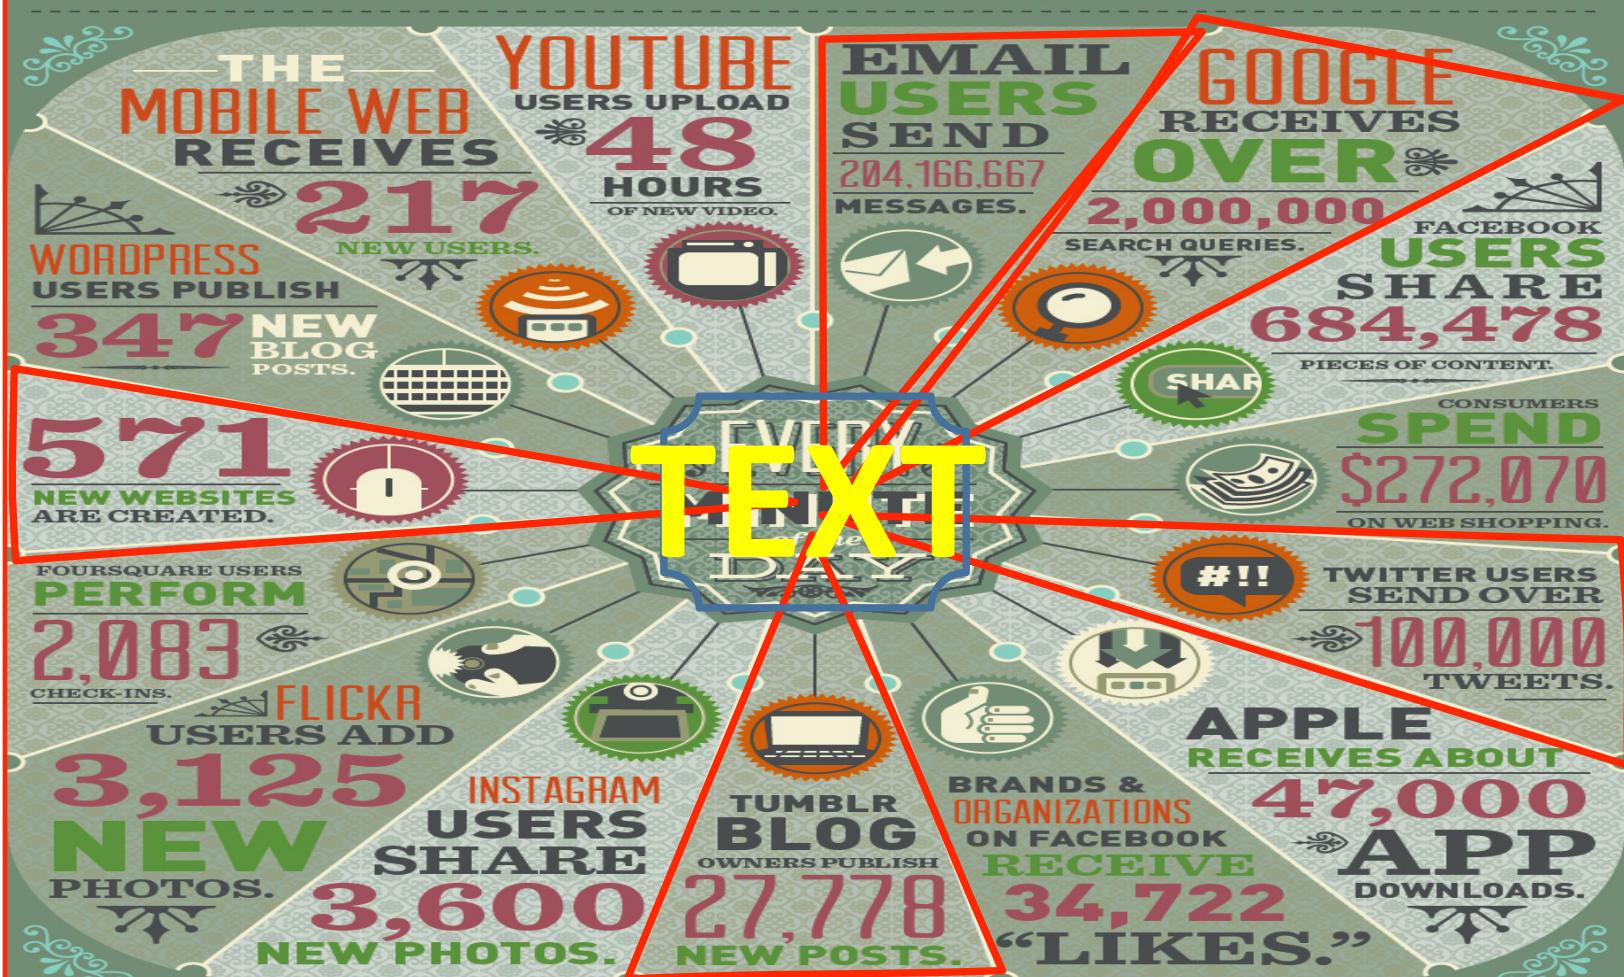
# Graph-of-words: boosting text mining with graphs



M. Vazirgiannis

Data Science & Mining group  
LIX @ Ecole Polytechnique

<http://www.lix.polytechnique.fr/dascim/>



# Text Mining - Where?

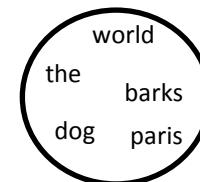
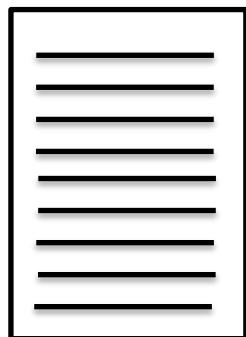
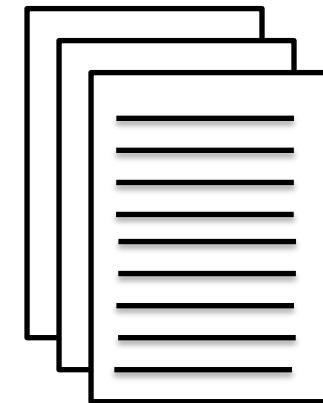
- Web search engines
  - Understand the user query (i.e. Jaguar...)
  - Find matching documents (ranking)
  - Group similar results together (clustering)
  - etc.
- Product recommendation
  - Understand product descriptions
  - Understand product reviews
  - etc.

# Text Mining

- Analogy with the petroleum industry:  
*Explore, extract, refine & use*
- Here *the oil is the information* contained in textual documents
- Text mining involves:
  - Data Mining & Machine Learning (ML)
  - Information Retrieval (IR)
  - Natural Language Processing (NLP)

# Text Mining – Terminology

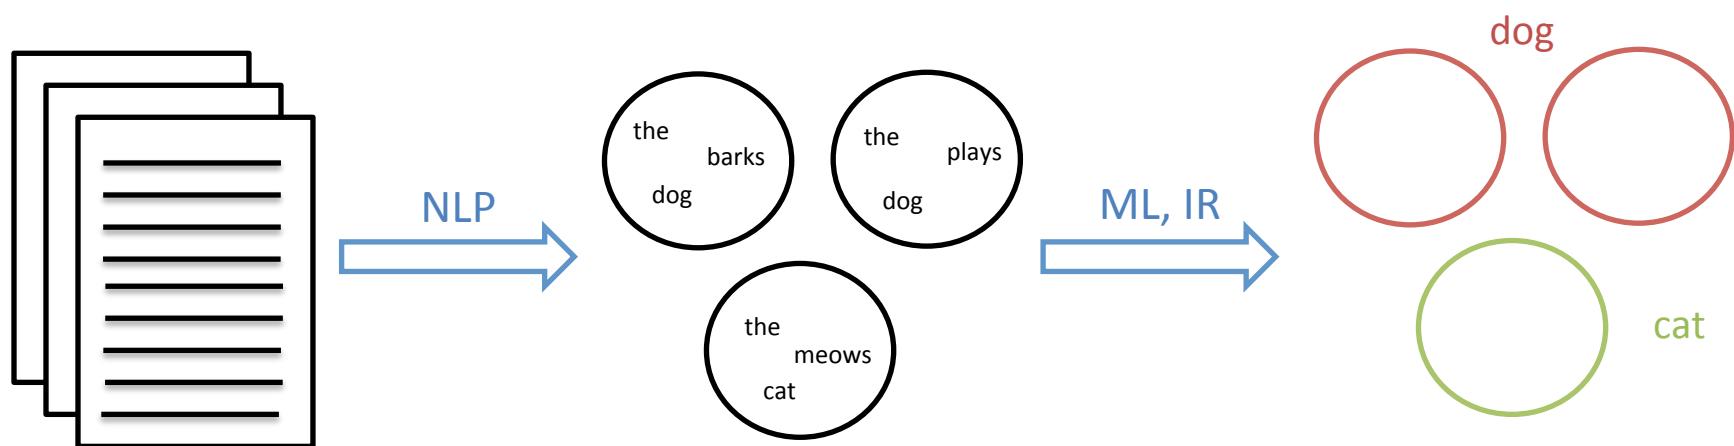
- You apply text mining on a collection of documents:
  - the collection is the data set
  - the documents are the data points
- Because text is unstructured, a document is usually converted in a common *representation*



bag-of-words

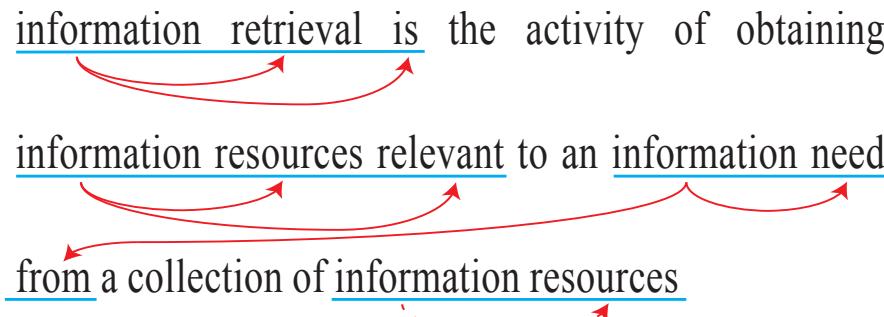
# Text Mining – Pipeline

- Input: raw text
- Output: some information (application-dependent)
- In-between:
  - 1) document representation
  - 2) feature extraction – **bag of words**



# Bag of Words - issues

information retrieval is the activity of obtaining  
information resources relevant to an information need  
from a collection of information resources



- term independence assumption
- term frequency weighting

Bag of words: ((activity,1), (collection,1)  
(information,4), (relevant,1),  
(resources, 2), (retrieval, 1)..)

# Graph of Words – a novel approach for text mining

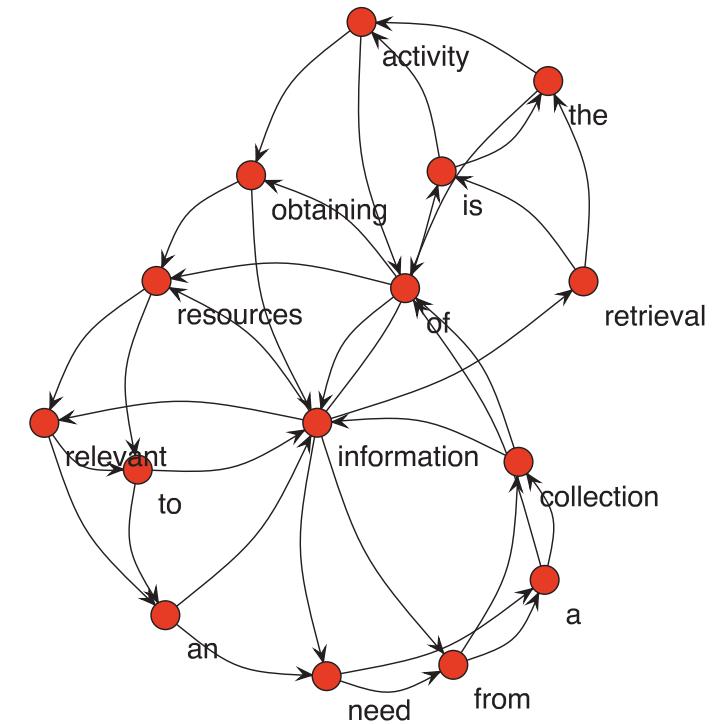
- Challenge “term independence” and “term frequency weighting” assumptions taking into account **word dependence, order** and **distance**
- Employ a graph-based document representation capturing the above
- Graphs successfully used in IR to encompass relations and propose meaningful weights (e.g. PageRank)

# Graph of Words (GoW)

information retrieval is the activity of obtaining  
information resources relevant to an information need  
from a collection of information resources

Bag of words: ((activity,1), (collection,1)  
(information,4), (relevant,1),  
(resources, 2), (retrieval, 1)..)

Captures: frequency, order  
and distance, ...



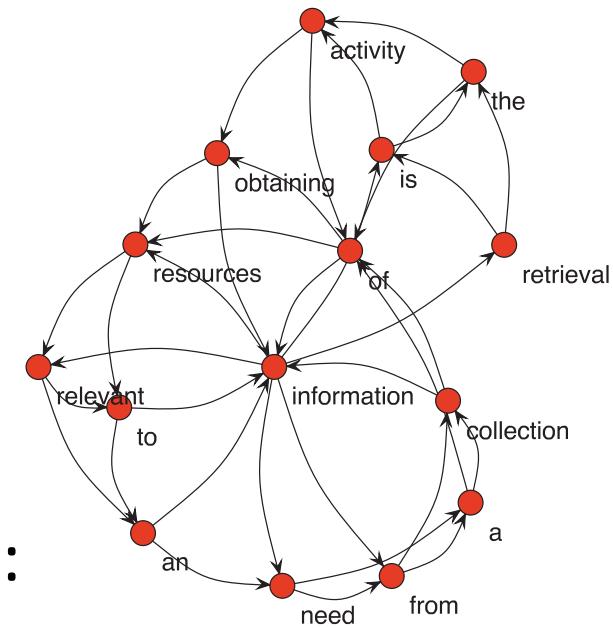
# GoW for text mining

- 1. Ad Hoc Information Retrieval
- 2. Single-Document Keyword Extraction
  - Elect *representative* words for a document that best describe it.
  - Sub-event Detection in Textual Streams (twitter)
  - keyword extraction Chrome extension
- 3. Text Categorization
  - as a Graph Classification Problem
  - Sentence based Kernels for short text categorization

# Graph-based Ad Hoc IR [CIDKM2013]

## Ad Hoc Information Retrieval [1,2]

- Unweighted directed graph
- word weighting in the document:  
number of neighbors in the graph
  - => favor words that co-occur with many different other words
- Robust to varying document length:
  - weight of a word increases only with **new context** of co-occurrences as opposed to the word frequency that increases with any new co-occurrence.



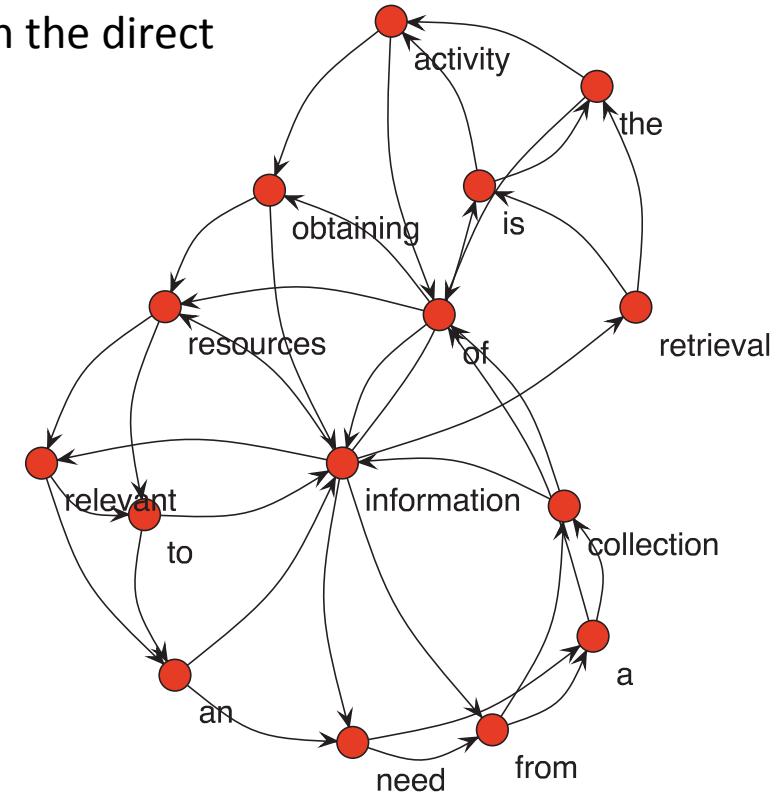
# Other efforts involving graphs

- Graph representations of text [Blanco and Lioma, 2012]:
  - To challenge the traditional bag-of-words document representation
- take into account some term dependence and term order
- node as linguistic unit (i.e., sentence, word or a character)
- edge as linguistic relationship (syntactic or semantic)
- syntactic graph-of-words with edges representing co-occurrences in a small sliding window.

# Indegree-BASED TW

- The weight of a term in a document is its **indegree** (numbers of incoming edges) in the **graph-of-word**
- It represents the **number of distinct contexts of occurrence**
- We store the document as a vector of weights in the direct index and similarly in the inverted index
- Example:

information	5
retrieval	1
is	2
the	2
activity	2
of	3
obtaining	2
resources	3
relevant	2
to	2
an	2
need	2
from	2
a	2
collection	2



# TF-IDF and BM25

Term  $t$ , document  $d$ , collection of size  $N$ , term frequency  $tf(t, d)$ , document frequency  $df(t)$ , document length  $|d|$ , average document length  $avdl$ , asymptotical marginal gain  $k_1$  (1.2), slope parameter  $b$

**TF-IDF** [Singhal et al., TREC-7]

$$\text{TF-IDF}(t, d) = \left( \frac{1 + \log(1 + \log(tf(t, d)))}{1 - b + b \times \frac{|d|}{avdl}} \right) \times \log\left(\frac{N+1}{df(t)}\right)$$

**BM25** [Lv and Zhai, CIKM '11]

$$\text{BM25}(t, d) = \left( \frac{(k_1 + 1) \times tf(t, d)}{k_1 \times \left(1 - b + b \times \frac{|d|}{avdl}\right) + tf(t, d)} \right) \times \log\left(\frac{N+1}{df(t)}\right)$$

# TW-IDF

- $tw$ : indegree of the vertex representing the term in the graph

$$\text{TW-IDF}(t, d) = \left( \frac{tw(t, d)}{1 - b + b \times \frac{|d|}{avdl}} \right) \times \log \left( \frac{N + 1}{df(t)} \right)$$

- Term  $t$ , document  $d$ , collection of size  $N$ , term weight  $tw(t, d)$ , document frequency  $df(t)$ , document length  $|d|$ , average document length  $avdl$ , slope parameter  $b$

# Experiments - Datasets

- **Disks 1 & 2 (TREC)**

741,856 news articles from Wall Street Journal (1987-1992), Federal Register (1988-1989), Associated Press (1988-1989) and Information from the Computer Select disks (1989-1990)

- **Disks 4 & 5 (TREC, minus the Congressional Record)**

528,155 news releases from Federal Register (1994), Financial Times (1991-1994), Foreign Broadcast Information Service (1996) and Los Angeles Times (1989-1990)

- **WT10G (TREC)**

1,692,096 crawled pages from a snapshot of the Web in 1997

- **.GOV2 (TREC)**

25,205,179 crawled Web pages from .gov sites in early 2004

# Datasets (cont.)

Statistic \ Dataset	Disks 1 & 2	Disks 4 & 5	WT10G	.GOV2
Statistic				
# of documents	741,856	528,155	1,692,096	25,205,179
# of unique terms	535,001	520,423	3,135,780	15,324,292
average # of terms ( <i>avdl</i> )	237	272	398	645
average # of vertices	125	157	165	185
average # of edges	608	734	901	1,185

**Table 1:** Statistics on the four TREC datasets used; Disks 4&5 excludes the Congressional Record.  
The average values are computed per document.

# EVALUATION

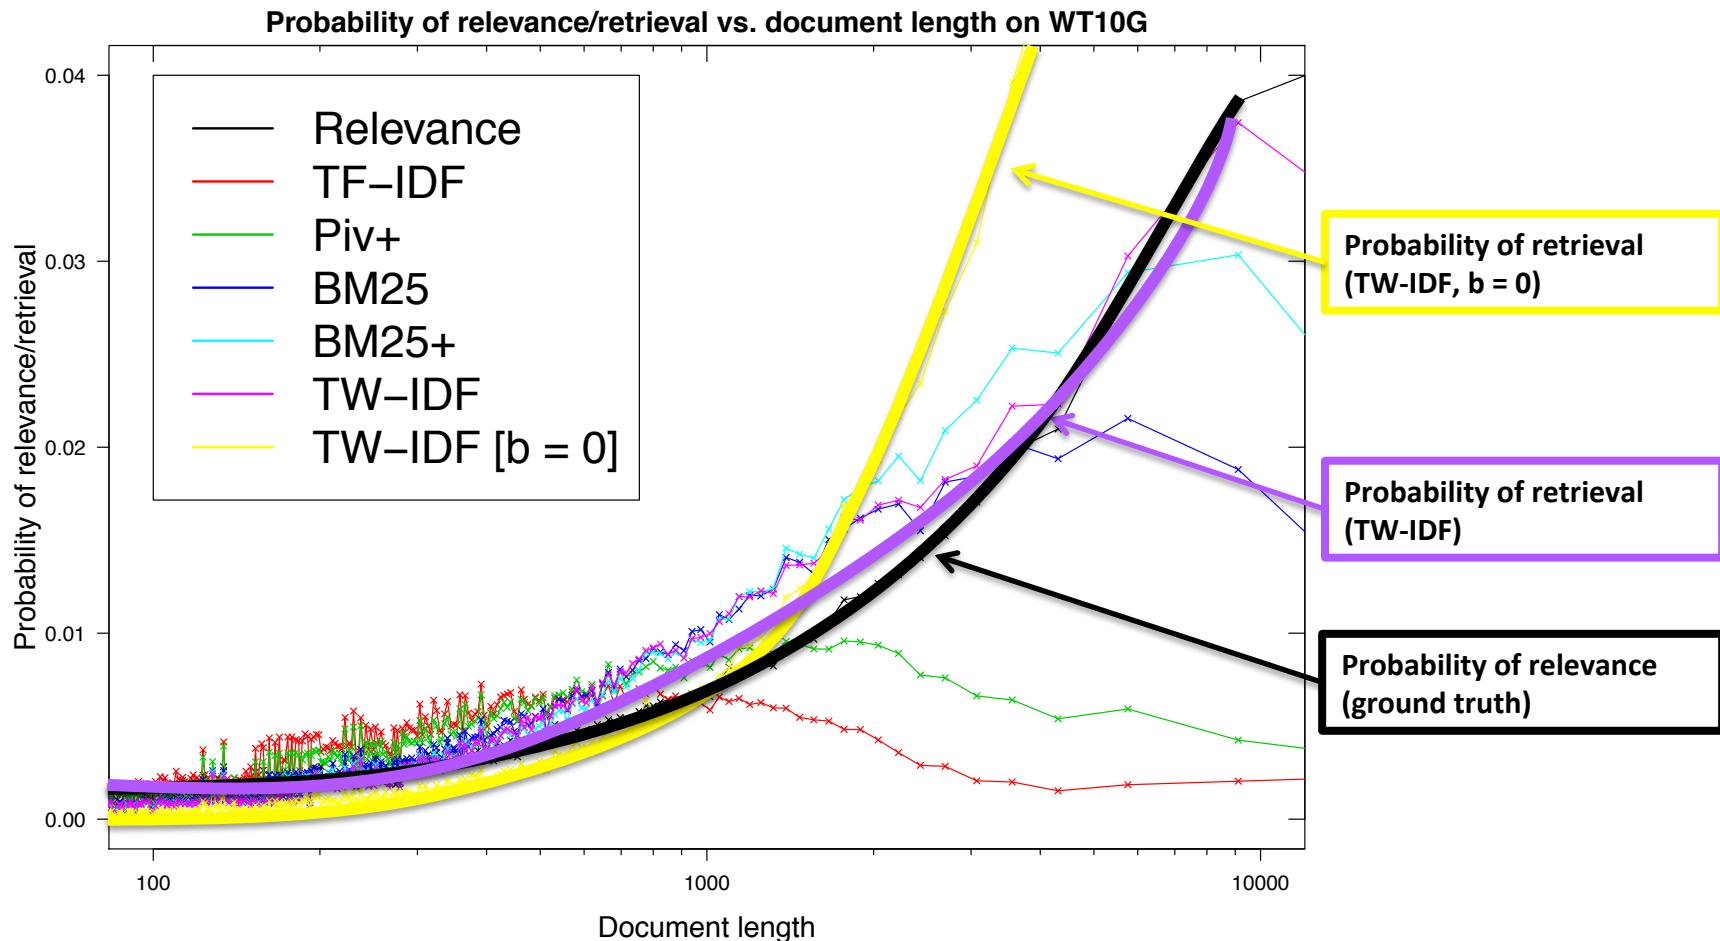
- Mean Average Precision (MAP) and Precision at 10 (P@10)  
considering only the top-ranked 1000 documents for each run
- Statistical significance of improvement was assessed using the Student's paired t-test
  - Two-sided p-values less than 0.05 and 0.01 to reject the null hypothesis
- Likelihood of relevance vs. likelihood of retrieval  
[Singhal et al., SIGIR '96]
- 4 baseline models: TF-IDF, BM25, Piv+ and BM25+

# Graph-based Ad Hoc IR II

- Evaluation in terms of:
  - Mean Average Precision
  - Precision@10
  - Probability of relevance  
vs. probability of retrieval

Model	$b$	TREC1-3 Ad Hoc		TREC 2004 Robust		TREC9-10 Web		TREC 2004-2006 Terabyte	
		MAP	P@10	MAP	P@10	MAP	P@10	MAP	P@10
TF <sub>pol</sub>	0.20	0.1471	0.3960	0.1797	0.3647	0.1260	0.1875	0.1853	0.4913
TF <sub>kop</sub>	0.75	0.1346	0.3533	0.2045	0.3863	0.1702	0.2208	0.2527	0.5342
TW	none	0.1502	0.3662	0.1809	0.3273	0.1430	0.1979	0.2081	0.5021
TW <sub>p</sub>	0.003	<b>0.1576**</b>	<b>0.4040**</b>	<b>0.2190**</b>	<b>0.4133**</b>	<b>0.1946**</b>	<b>0.2479**</b>	<b>0.2828**</b>	<b>0.5407**</b>
TF-IDF	0.20	0.1832	0.4107	0.2132	0.4064	0.1430	0.2271	0.2068	0.4973
BM25	0.75	0.1660	0.3700	0.2368	0.4161	0.1870	0.2479	0.2738	0.5383
TW-IDF	0.003	<b>0.1973**</b>	<b>0.4148*</b>	<b>0.2403**</b>	<b>0.4180*</b>	<b>0.2125**</b>	<b>0.2917**</b>	<b>0.3063**</b>	<b>0.5633**</b>

# Likelihood of relevance vs. likelihood of retrieval



# GoW for text mining

- 1. Ad Hoc Information Retrieval
- 2. Single-Document Keyword Extraction
  - Elect *representative* words for a document that best describe it.
  - Sub-event Detection in Textual Streams (twitter)
  - keyword extraction Chrome extension
- 3. Text Categorization
  - as a Graph Classification Problem
  - Sentence based Kernels for short text categorization

# Single Document Keyword Extraction

## [ECIR 2015]

Keywords are used everywhere:

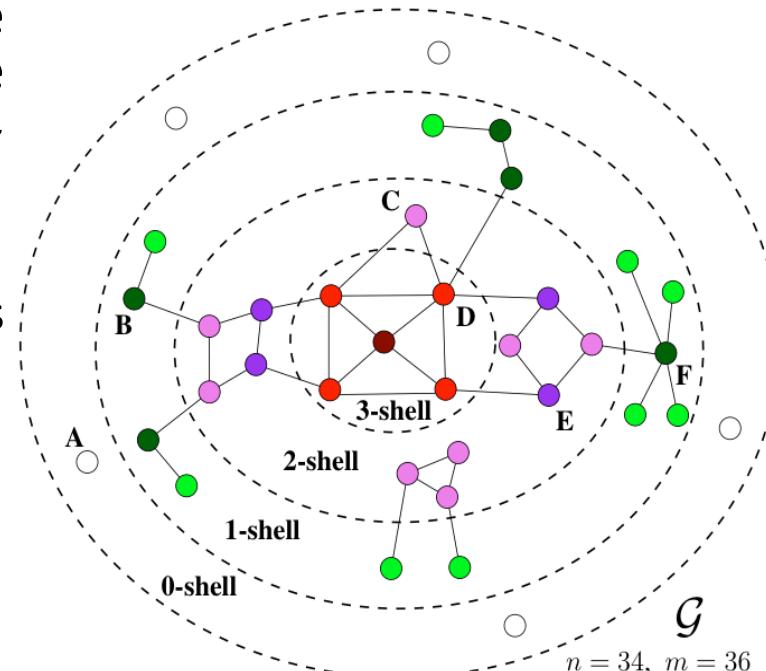
- looking up information on the Web (e. g., via a search engine bar)
- finding similar posts on a blog (e. g., tag cloud)
- for ads matching (e. g., AdWords' keyword planner)
- for research paper indexing and retrieval (e. g., SpringerLink)
- for research paper reviewer assignment

Applications are numerous:

- **summarization** (to get a gist of the content of a document)
- **information filtering** (to select specific documents of interest)
- **indexing** (to answer keyword-based queries)
- **query expansion** (using additional keywords from top results)

# Graph degeneracy – k-core

- Let  $k$  be an integer.
- A subgraph  $H_k = (V', E')$ , induced by the subset of vertices  $V' \subseteq V$  (and the subset of edges  $E' \subseteq E$ ), is called a  $k$ -core or *core of order k* iff:
  - $\forall v \in V'$ ,  $\deg_{H_k}(v) \geq k$  and
  - $H_k$  is the maximal subgraph with this property, i.e. it cannot be augmented without losing this property.
- $k$ -core of a graph: maximal connected subgraph whose each vertex has at least degree  $k$  within the subgraph.
- The core of maximum order is called the *main core*.

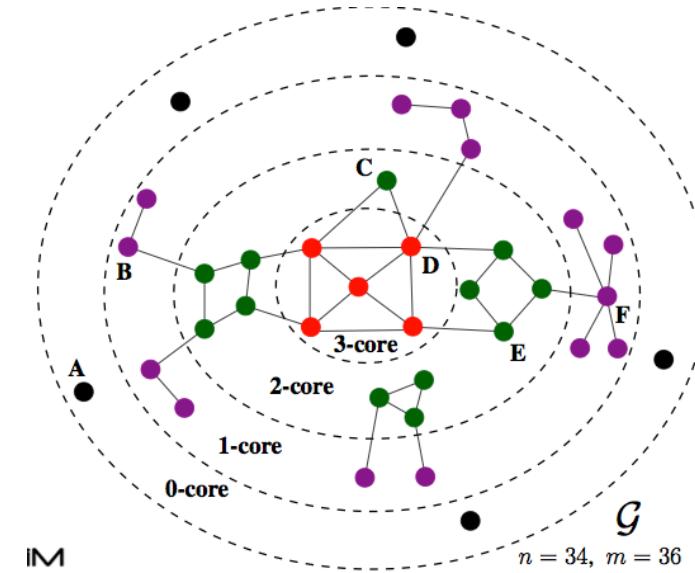


# Graph-based keyword extraction

Existing graph-based keyword extractors:

- PageRank [Mihalcea and Tarau, 2004]
- HITS [Litvak and Last, 2008]
- assign a *centrality* based influence score to a node
- top ones the most representative

⇒ we propose to retain the main core of the graph to extract the nodes based on their centrality and cohesiveness.



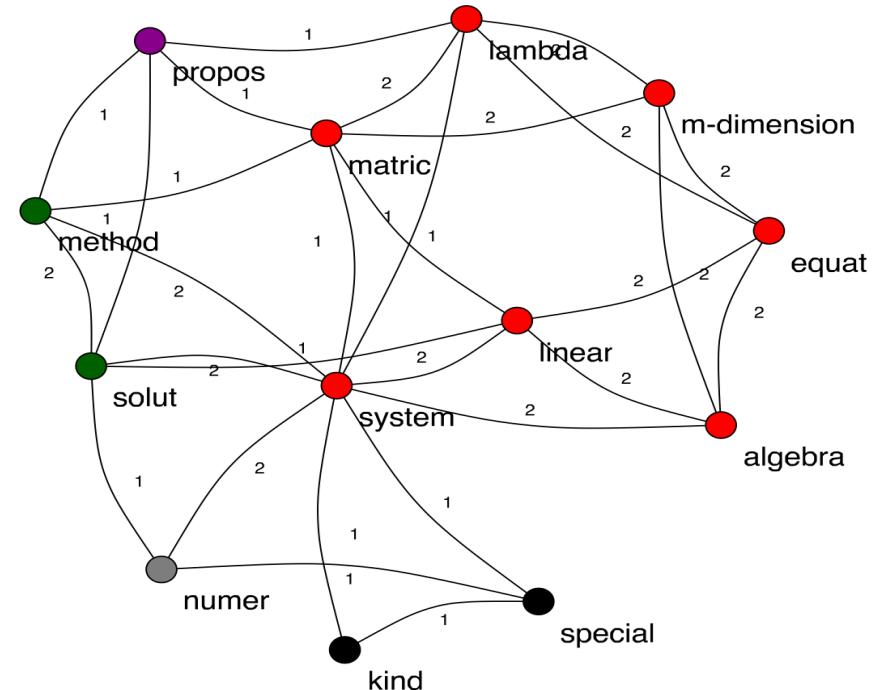
K-core decomposition of the graph

# Graph-based Keyword Extraction

- Single-Document Keyword Extraction

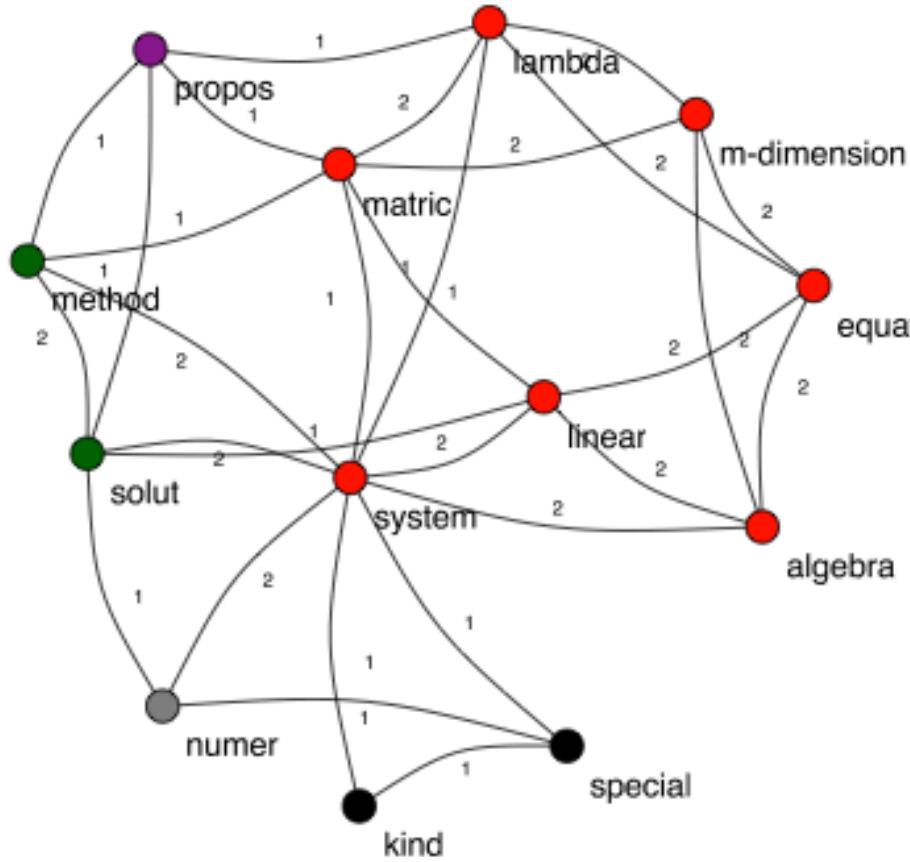
- Elect the most cohesive sets of words in the graph as keywords
- Use k-core decomposition to extract the main core of the graph
- Weighted edges as opposed to Ad Hoc IR (single-document => no normalization)

A method for solution of systems of linear algebraic equations with m-dimensional lambda matrices.  
A system of linear algebraic equations with m-dimensional lambda matrices is considered. The proposed method of searching for the solution of this system lies in reducing it to a numerical system of a special kind.



**Keywords manually assigned by human annotators**  
linear algebra equat; numer system; m-dimension lambda matric

# Pagerank vs. main core



Keywords manually assigned by human annotators  
 linear algebra equat; numer system; m-dimension lambda matric

WK-core	PageRank	WK-core	PageRank
<b>system</b>	6	<b>system</b>	1.93
<b>matric</b>	6	<b>matric</b>	1.27
<b>lambda</b>	6	<b>solut</b>	1.10
<b>linear</b>	6	<b>lambda</b>	1.08
<b>equat</b>	6	<b>linear</b>	1.08
<b>algebra</b>	6	<b>equat</b>	0.90
<b>m-dim...</b>	6	<b>algebra</b>	0.90
<b>method</b>	5	<b>m-dim...</b>	0.90
<b>solut</b>	5	<b>propos</b>	0.89
<b>propos</b>	4	<b>method</b>	0.88
<b>numer</b>	3	<b>special</b>	0.78
<b>specia</b>	2	<b>numer</b>	0.74
<b>kind</b>	2	<b>kind</b>	0.55

# Keywords are not unigrams

- 500 abstracts from the *Inspec* database used in our experiments,
- 4,913 keywords manually assigned by human annotators
- only 662 are unigrams (13%).
- Bigrams (2,587 – 52%) ... 7-grams (5).
  - ⇒ keywords are bigrams, if not higher order n-grams.
  - ⇒ the interactions within keywords need to be captured in the first place – i.e. in the graph.
  - ⇒ we can consider a k-core to form a “long-distance (k +1)-gram” [Bassiou and Kotropoulos, 2010]

# K-cores are adaptive

- Most techniques in keyword extraction assign a score to each feature and then take the top ones.
  - But how many? Absolute number (top X) or relative number (top X%)?
  - Besides, at fixed document length, humans may assign more keywords for a document than for another one.
- ⇒ X be decided at document-level (*size of the main core*).

# Data sets

- *Hulth2003* – 500 abstracts from the *Inspec* database [Hulth, 2003]
- *Krapi2009* – 2,304 ACM full papers in Computer Science (references and captions excluded) [Krapivin et al., 2009]

All approaches are *unsupervised* and single-document

# Models

Graph-of-words:

- undirected edges
- forward edges (natural flow of the text – an edge **term1 → term2** meaning that **term1** precedes **term2** in a sliding window)
- backward edges

Keyword extractors:

- PageRank
  - HITS (authority scores only)
  - K-core
  - Weighted K-core
- } Top 33% or top 15% keywords
- } Main core

# Evaluation metrics

- Each document has a set of golden keywords assigned by humans.
- ⇒ **precision, recall** and **F1-score** per document
- ⇒ **macro-average** each metric at the collection level

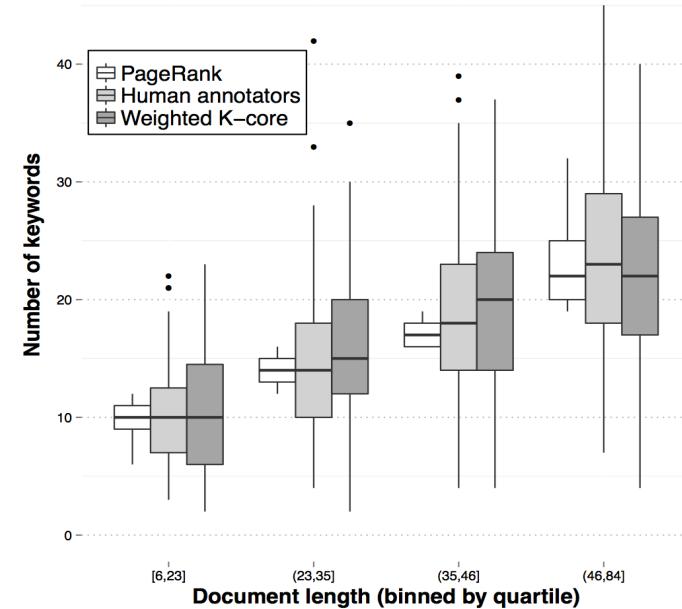
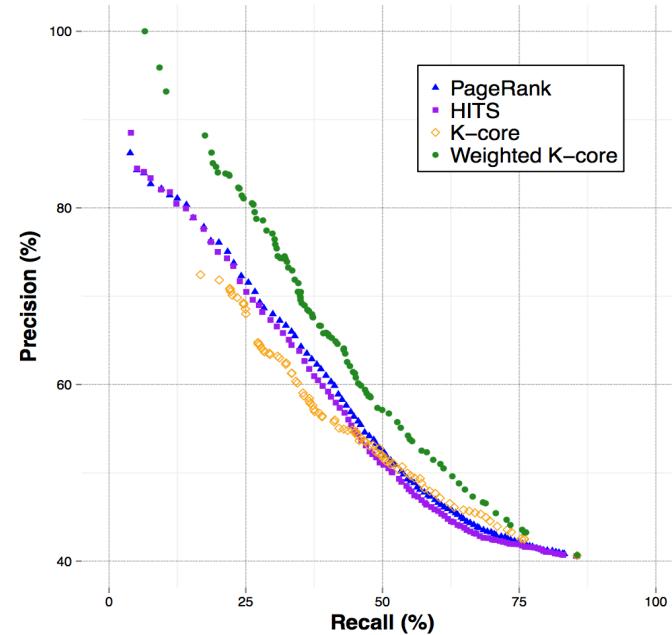
# Evaluation metrics 2/2

- Given a golden bigram to extract, how do you penalize a system that retrieves part of it (unigram) or more than it (trigram)?
- ⇒ golden keywords are transformed into unigrams for standard precision and recall definitions.
- ⇒ “reconciliation” step considered as a post-processing step.

# Graph-based Keyword Extraction – Experimental results

- Evaluation in terms of:

- Precision
- Recall
- F1-score
- Precision/recall
- # of keywords



Graph	Dataset	Macro-averaged precision (%)				Macro-averaged recall (%)				Macro-averaged F1-score (%)			
		PageRank	HITS	K-core	WK-core	PageRank	HITS	K-core	WK-core	PageRank	HITS	K-core	WK-core
undirected edges	Hulth2003	58.94	57.86	46.52	<b>61.24*</b>	42.19	41.80	<b>62.51*</b>	50.32*	47.32	46.62	49.06*	<b>51.92*</b>
	Krapi2009	50.23	49.47	40.46	<b>53.47*</b>	48.78	47.85	<b>78.36*</b>	50.21	49.59	47.96	46.61	<b>50.77*</b>
forward edges	Hulth2003	55.80	54.75	42.45	<b>56.99*</b>	41.98	40.43	<b>72.87*</b>	46.93*	45.70	45.03	<b>51.65*</b>	50.59*
	Krapi2009	47.78	47.03	39.82	<b>52.19*</b>	44.91	44.19	<b>79.06*</b>	45.67	45.72	44.95	46.03	<b>47.01*</b>
backward edges	Hulth2003	59.27	56.41	40.89	<b>60.24*</b>	42.67	40.66	<b>70.57*</b>	49.91*	47.57	45.37	45.20	<b>50.03*</b>
	Krapi2009	51.43	49.11	39.17	<b>52.14*</b>	49.96	47.00	<b>77.60*</b>	50.16	<b>50.51</b>	47.38	46.93	50.42

# Conclusions – lessons learned

- Explored single-document keyword extraction using k-core on graph-of-words:
  - Capitalized on graph representations of text to extract central terms.
  - Captured more cohesive subgraphs of words central and densely connected.
  - Extracted keywords are more likely to form higher order n-grams and their number adapts to the graph structure.

# Eventual example

- Stemmed unigrams of the main core of the graph- of-words of the paper document:  
*{keyword, extract, graph, represent, text, weight, graph-of-word, k-core, degeneraci, edg, vertic, number, document}*.
- Using PageRank, “**work**” appears in the top 5, “**term**” and “**pagerank**” in the top 10, and “**case**” and “**order**” in the top 15. Central words but not in cohesion with the rest and probably not relevant.

# GoW for text mining

- 1. Ad Hoc Information Retrieval
- 2. Single-Document Keyword Extraction
  - Elect *representative* words for a document that best describe it.
  - Sub-event Detection in Textual Streams (twitter)
  - keyword extraction Chrome extension
- 3. Text Categorization
  - as a Graph Classification Problem
  - Sentence based Kernels for short text categorization

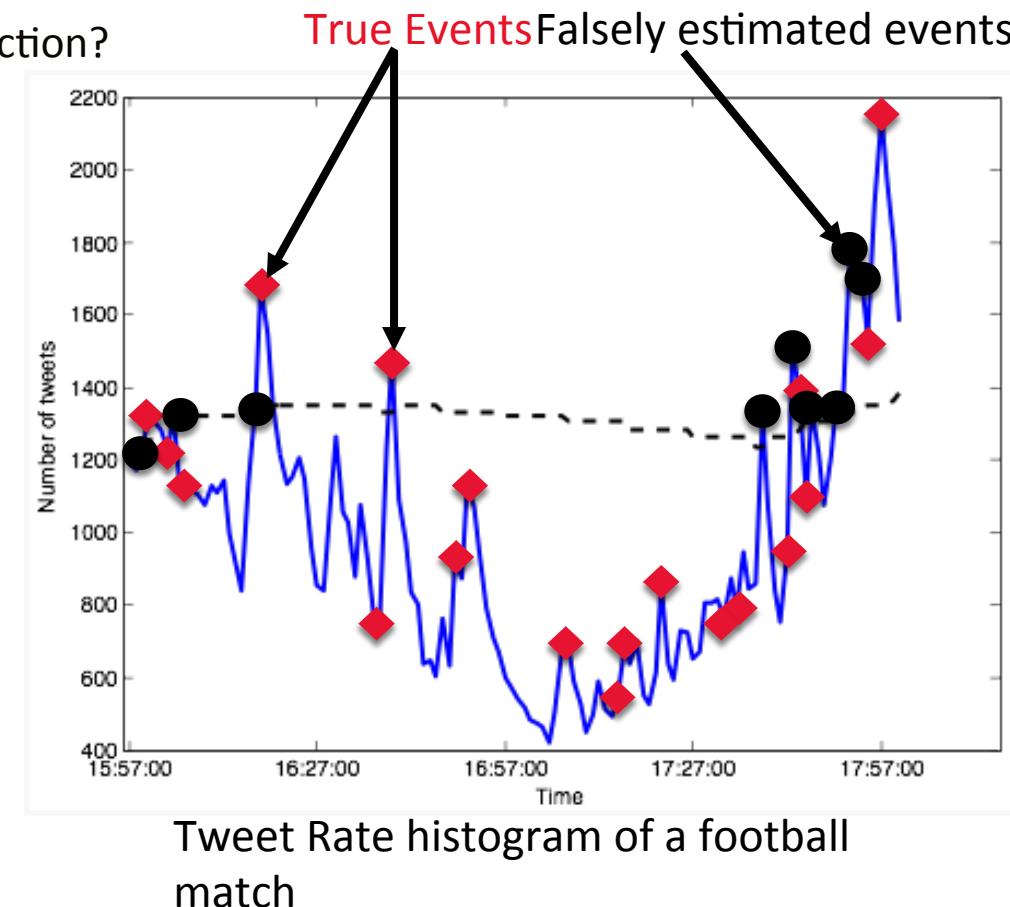
# Sub-event Detection in Textual Streams (twitter)

## [AAAI/ICWSM 2015]

1. Large volume of documents in social media
2. Events are not covered by traditional media
3. News appear fast in Twitter
4. Is Tweet Rate suited for sub-event Detection?

### Contribution

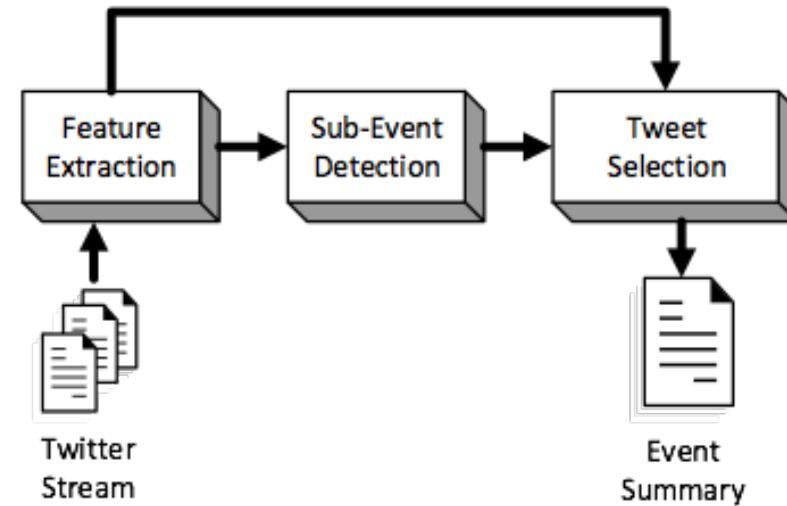
*A novel real-time detection mechanism that accurately detects sub-events in an unsupervised and out-of-the-box manner.*



# System Architecture

# Real Time Event Summarization

1. Feature Extraction: Extracts the terms that best describe the current state of the event
  2. Sub-Event Detection: Decides whether a sub-event has occurred
  3. Tweet Selection: Ranks all the tweets and selects the first one.



# System Architecture

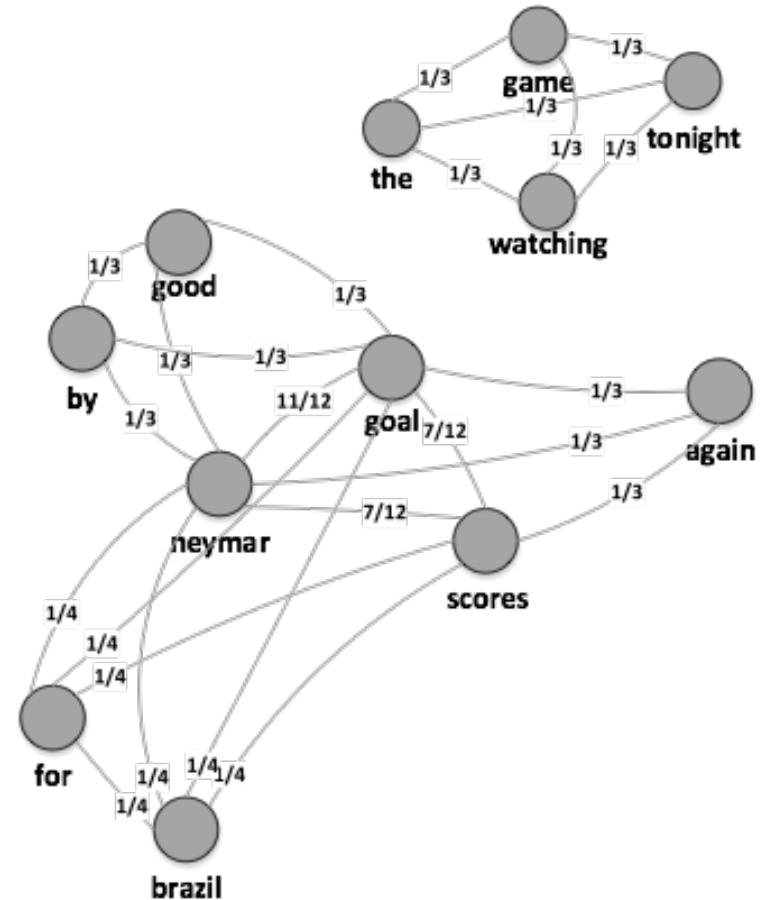
- Steps are repeated every 60 seconds
  - The summary of the whole event is constructed by aggregating the individual sub-event descriptions

# Graph-of-Words from Twitter Streams

- Represents all the input tweets  
[Rousseau and Vazirgiannis, CIKM 2013]
- node  $\leftrightarrow$  unique term
- edge  $\leftrightarrow$  #co-occurrences  
within a tweet (normalized)

Example Graph:

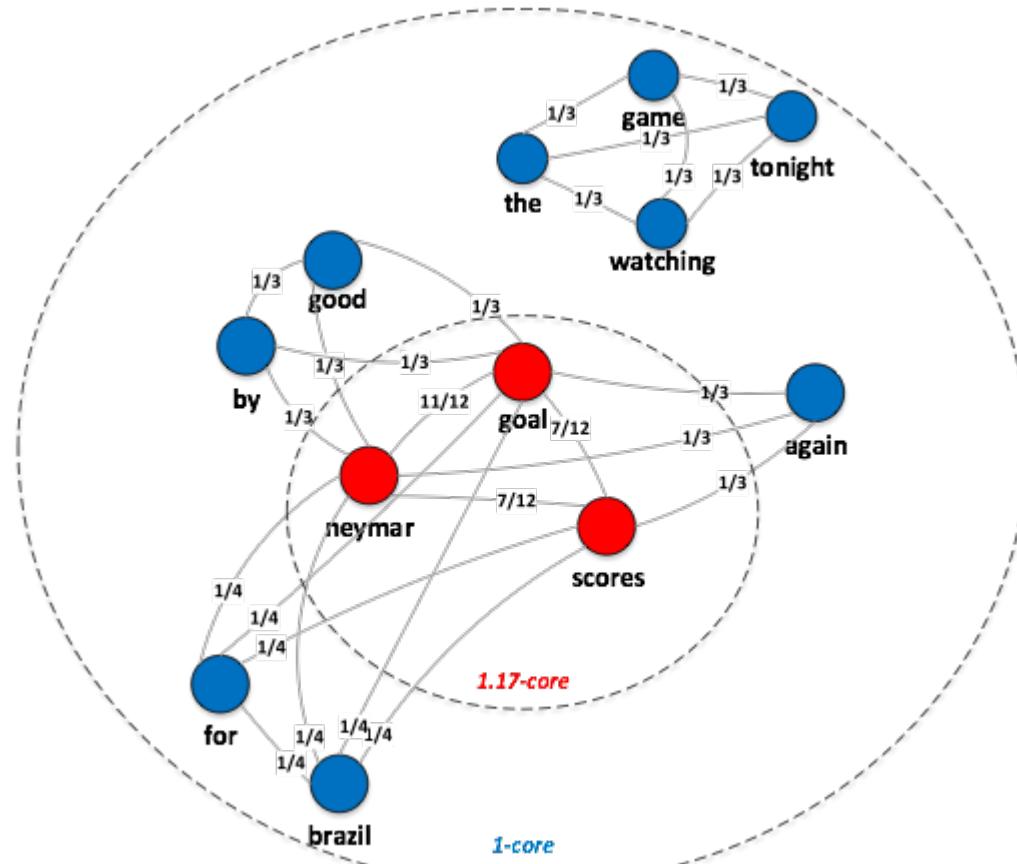
- 1) Good goal by Neymar
- 2) Goal! Neymar scores for brazil
- 3) Goal!! Neymar scores again
- 4) Watching the game tonight



A graph-of-words built from 4 tweets

# Graph Degeneracy for Feature Extraction

- Each term is given a score corresponding to its core number
- The  $k$ -core of a graph is the maximal subgraph whose vertices are at least of degree  $k$  within the subgraph.
- A vertex has index  $k$  if it belongs to the  $k$ -core but not to the  $k+1$  core



k-core decomposition of the Graph-of-Words

# Sub-event Detection

$$\sum_{i=1}^d c_i^t > \theta \times \frac{1}{p} \sum_{j=t-p}^{t-1} \sum_{i=1}^d c_i^j$$

$c_i^t$  Core number of term  $i$  at time slot  $t$

$d$  Number of terms selected

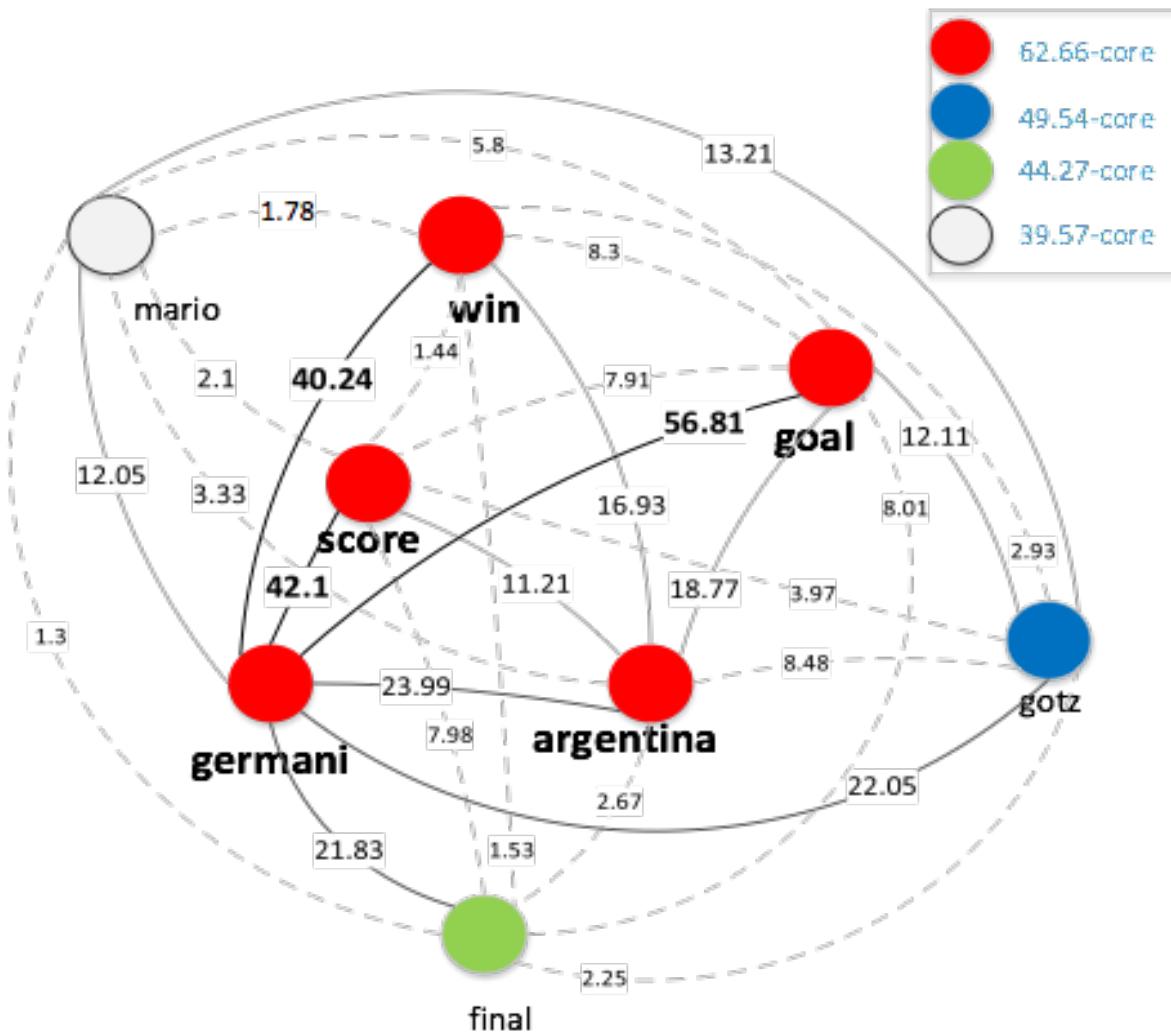
$\theta$  Decision Threshold

$p$  Number of previous time slots

Sub-event Detection steps:  
(Every 60 seconds)

- 1) Extract the top  $d$  terms with highest weights
- 2) Sum the term weights
- 3) If it exceeds the threshold a sub-event is detected

# Germany's goal, 2014 World Cup: Graph-of-Words

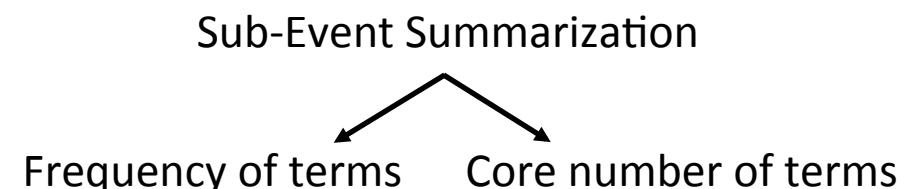
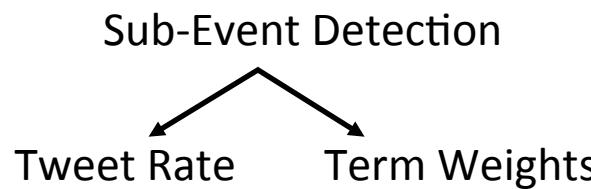


Snapshot of the four highest cores of the graph generated after Germany's goal in the 2014 FIFA World Cup final

# Tweet Selection as Sub-event Summarization

- Activated only if a sub-event has been detected
- Tweets are scored based on the *sum of their term weights*
- Selects the most informative tweet of the sub-event
  - The tweet with the highest score is chosen

# Experimental Setup



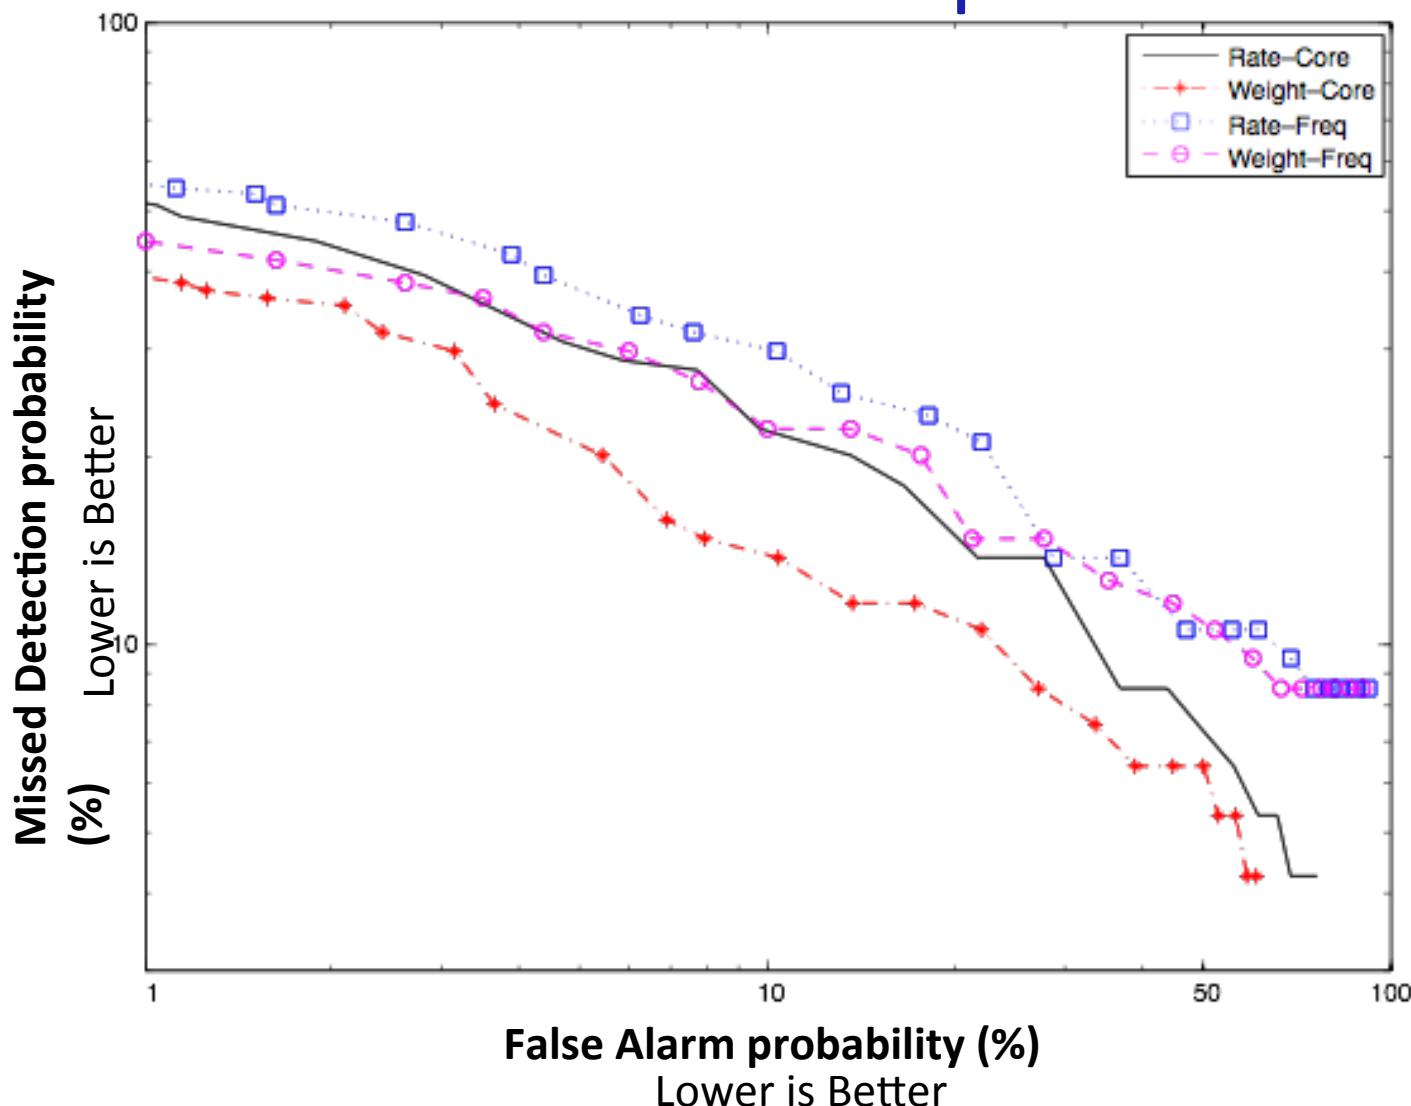
- Baselines-Approaches  
(Detection -Term Weight)
- Rate–Freq: the common baseline
  - Rate–Core
  - **Weight–Core**: Our approach
  - Weight–Freq

# Dataset

Match	#sub-events	#tweets
Germany - Argentina	8	1,907,999
Argentina - Belgium	7	1,355,472
France - Germany	6	1,321,781
Honduras-Switzerland	7	168,519
Greece - Ivory Coast	10	251,420
Croatia - Mexico	11	600,776
Cameroon - Brazil	11	532,756
Netherlands - Chile	7	301,067
Australia - Spain	9	252,086
Germany - Ghana	8	718,709
Australia - Netherlands	11	126,971
<b>All Matches</b>	<b>95</b>	<b>7,537,556</b>

FIFA 2014 World Cup Dataset

# Sub-event Detection performance



Method	Micro F1-score	Macro F1-score
<i>Weight-Core</i>	<b>0.68</b>	<b>0.72</b>
Rate-Core	0.61	0.63
Weight-Freq	0.61	0.64
Rate-Freq	0.54	0.60

Average micro and macro **F1-score** over 11 matches for the 4 considered approaches.

Event type	#actual Events	#detected Events
Goal	32	30
Penalty	2	2
Red Card	1	0
Yellow Card	27	14
Match Start	11	8
Match End	11	11
Half Time	11	10

Number of sub-events Detected

# Tweet Summarization performance

Time	Our Summary	ESPN FC
8'	Goal!!!!Argentina!! After eight minutes Argentina lead Belgium by 1-0 scored by Higuain	Goal! Argentina 1, Belgium 0. Gonzalo Higuain (Argentina) right footed shot from the centre of the box to the bottom left corner.
45'+2'	HT: Argentina 1-0 Belgium. Fantastic goal by Higuain gives Argentina the slight lead over the red devils.	First Half ends, Argentina 1, Belgium 0.
52'	52m - Belgium's Eden Hazard with the first yellow card of the game	Eden Hazard (Belgium) is shown the yellow card for a bad foul.
75'	Argentina 1 - 0 Belgium   Biglia booked a yellow card. Meanwhile, Chadli on for Eden Hazard.	Lucas Biglia (Argentina) is shown the yellow card for a bad foul.
90+5'	Well at least that goal makes them advance to the semi finals. Argentina gets the ticket to advance and Belgium goes home.	Match ends, Argentina 1, Belgium 0.

Summary of the Argentina vs. Belgium match generated automatically using Weight–Core and manually by ESPN.

# Conclusion

- We proposed a sub-event detection approach based on the k-core decomposition on graph-of-words
- The algorithm exploits the fact that the vocabulary of tweets gets more specific when a sub-event occurs
- Our detection mechanism is able to accurately detect important moments as they occur
- The tweets selected by our system give an overview of the event

# GoW for text mining

- 1. Ad Hoc Information Retrieval
- 2. Single-Document Keyword Extraction
  - Elect *representative* words for a document that best describe it.
  - Sub-event Detection in Textual Streams (twitter)
  - keyword extraction Chrome extension
- 3. Text Categorization
  - as a Graph Classification Problem
  - Sentence based Kernels for short text categorization

# Chrome addon for keyword extraction

- Based on the main core of the document [ECIR 2015]
- See [demo](#)

# GoW for text mining

- 1. Ad Hoc Information Retrieval
- 2. Single-Document Keyword Extraction
  - Elect *representative* words for a document that best describe it.
  - Sub-event Detection in Textual Streams (twitter)
  - keyword extraction Chrome extension
- 3. Text Categorization
  - as a Graph Classification Problem
  - Sentence based Kernels for short text categorization

# Text Categorization as a Graph Classification Problem [ACL2015]

- Single-label multi-class text categorization
- **Graph-of-words** representation of textual documents
- Mining of frequent **subgraphs as features** for classification
- Main core retention to reduce the graph's sizes
- **Long-distance n-grams** more discriminative than standard n-grams

# Context

Applications of text classification are numerous:

- news filtering
- document organization
- spam detection
- opinion mining

Text documents classification compared to other domains:

- high number of features
- sparse feature vectors
- multi-class scenario
- skewed class distribution

# Background

Text categorization [Sebastiani, 2002, Aggarwal and Zhai, 2012]

- Standard baseline: unsupervised n-gram feature mining + supervised linear SVM learning
- Common approach for spam detection: same with Naive Bayes

⇒ n-grams to take into account some **word order** and some **word dependence** as opposed to unigrams

⇒ word inversion? subset matching?

# Background

- Graph classification
    - subgraphs as features
    - graph kernels [Vishwanathan et al., 2010]
  - Frequent subgraph feature mining
    - gSpan [Yan and Han, 2002]
    - FFSM [Huan et al., 2003]
    - Gaston [Nijssen and Kok, 2004]
  - expensive to mine all subgraphs, especially for “large” collections of “large” graphs
- ⇒ unsupervised discriminative feature selection?

# Subgraph-of-words

- A subgraph of size n corresponds to a long-distance n-gram  
⇒ takes into account **word inversion** and **subset matching**
- For instance, on the R8 dataset, {bank, base, rate} was a discriminative (top 5% SVM features) long-distance 3-gram for the category “interest”
  - “barclays **bank** cut its **base** lending **rate**”
  - “midland **bank** matches its **base rate**”
  - “**base rate** of natwest **bank** dropped”

!! patterns hard to capture with traditional n-gram bag-of-words.

# Graph of Words Classification

## Unsupervised feature mining and support selection

- gSpan mines the most frequent “subgraph-of-words” in the collection of graph-of-words
- subgraph frequency == long-distance n-gram document frequency
- minimum document frequency controlled via a **support** parameter
- the lower the support, the more features but the longer the mining, the feature vector generation and the learning
- ⇒ unsupervised support selection using the **elbow method** (inspired from selecting the number of clusters in k-means)

# Multiclass Scenario

- Text categorization ==  
multiple classes + skewed class distribution + single  
overall support value (local frequency)
- 100k features for majority classes vs. 100  
features for minority ones
- ⇒ mining per class with same relative support  
value

# Main core mining and n-gram feature selection

- Complexity to extract all features!  $\Rightarrow$  reduce the size of the graphs
- Maintain word dependence and subset matching  $\Rightarrow$  keep the densest subgraphs
  - $\Rightarrow$  retain the main core of each graph-of-words
  - use gSpan to mine frequent subgraphs in main cores
  - $\Rightarrow$  extract n-gram features on remaining text (terms in main cores)

# Experimental evaluation

## Standard datasets:

- *WebKB*: 4 most frequent categories among labeled webpages from various CS departments – split into 2,803 for training and 1,396 for test [Cardoso-Cachopo, 2007].
- *R8*: 8 most frequent categories of Reuters-21578, a set of labeled news articles from the 1987 Reuters newswire – split into 5,485 for training and 2,189 for test [Debole and Sebastiani, 2005].
- *LingSpam*: 2,893 emails classified as spam or legitimate messages – split into 10 sets for 10-fold cross validation [Androultsopoulos et al., 2000].
- *Amazon*: 8,000 product reviews over four different sub-collections (books, DVDs, electronics and kitchen appliances) classified as positive or negative – split into 1,600 for training and 400 for test each [Blitzer et al., 2007].

# Models

- 3 baseline models (n-gram features)
  - kNN (k=5)
  - Multinomial Naive Bayes (similar results with Bernoulli)
  - linear SVM
- 3 proposed approaches
  - gSpan + SVM (long-distance n-gram features)
  - MC + gSpan + SVM (long-distance n-gram features)
  - MC + SVM (n-gram features)

# Evaluation metrics

- Micro-averaged F1-score (accuracy, overall effectiveness)
- Macro-averaged F1-score (weight each class uniformly)
- Statistical significance of improvement in accuracy over the n-gram SVM baseline assessed using the micro sign test ( $p < 0.05$ )
- For the Amazon dataset, we report the average of each metric over the four sub-collections

# Effectiveness results (1/2)

Method	Dataset	WebKB		R8	
		Accuracy	F1-score	Accuracy	F1-score
kNN (k=5)		0.679	0.617	0.894	0.705
NB (Multinomial)		0.866	0.861	0.934	0.839
linear SVM		0.889	0.871	0.947	0.858
gSpan + SVM		<b>0.912*</b>	<b>0.882</b>	<b>0.955*</b>	<b>0.864</b>
MC + gSpan + SVM		0.901*	0.871	0.949*	0.858
MC + SVM		0.872	0.863	0.937	0.849

**Table:** Test accuracy and macro-average F1-score. Bold font marks the best performance in a column. \* indicates statistical significance at  $p < 0.05$  using micro sign test with regards to the SVM baseline of the same column. gSpan mining support values are 1.6% (WebKB) and 7% (R8).

# Effectiveness results (2/2)

Method	Dataset		LingSpam		Amazon	
	Accuracy	F1-score	Accuracy	F1-score		
kNN (k=5)	0.910	0.774	0.512	0.644		
NB (Multinomial)	0.990	0.971	0.768	0.767		
linear SVM	<b>0.991</b>	<b>0.973</b>	0.792	0.790		
gSpan + SVM	<b>0.991</b>	0.972	0.798*	0.795		
MC + gSpan + SVM	0.990	<b>0.973</b>	<b>0.800*</b>	<b>0.798</b>		
MC + SVM	0.990	0.972	0.786	0.774		

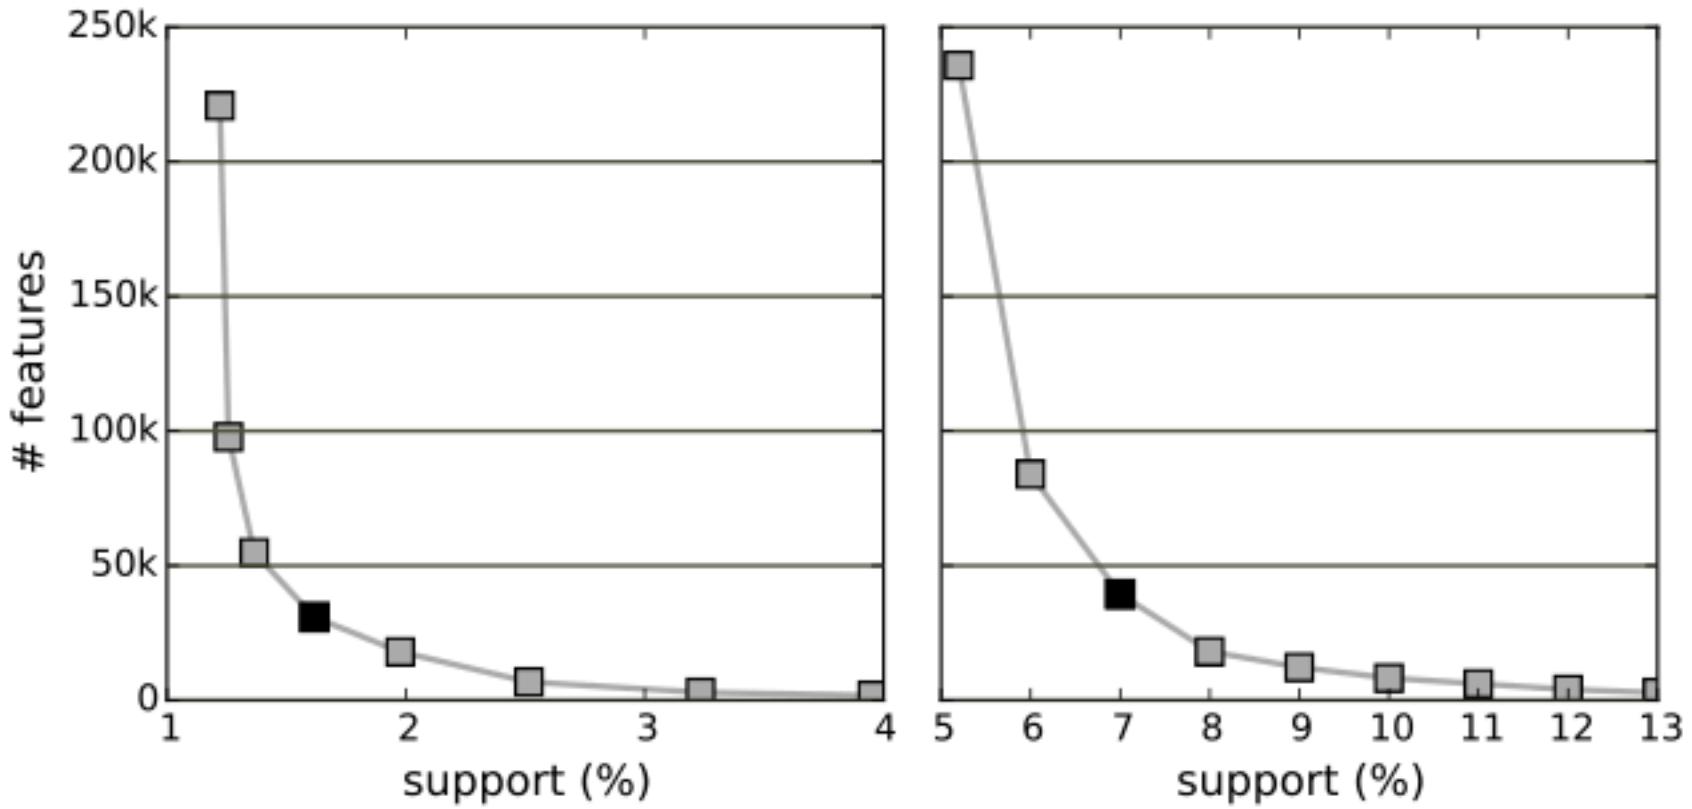
**Table:** Test accuracy and macro-average F1-score. Bold font marks the best performance in a column. \* indicates statistical significance at  $p < 0.05$  using micro sign test with regards to the SVM baseline of the same column. gSpan mining support values are 4% (LingSpam) and 0.5% (Amazon).

# Dimension reduction – main core

Dataset	# n-grams before	# n-grams after	reduction
WebKB	1,849,848	735,447	60 %
R8	1,604,280	788,465	51 %
LingSpam	2,733,043	1,016,061	63 %
Amazon	583,457	376,664	35 %

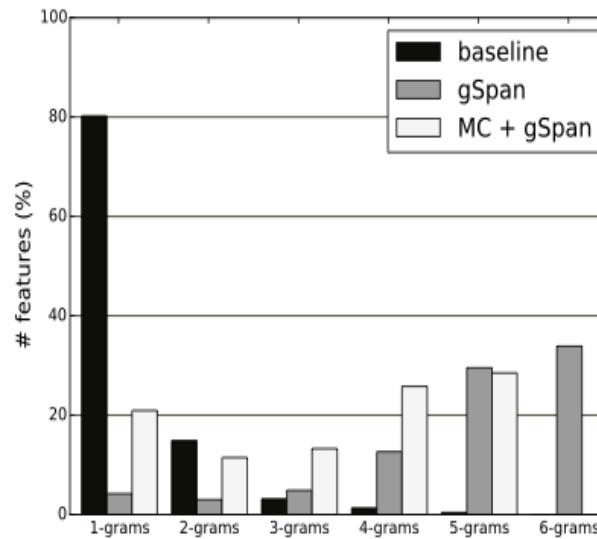
**Table:** Total number of n-gram features vs. number of n-gram features present only in main cores along with the reduction of the dimension of the feature space on all four datasets.

# Unsupervised support selection

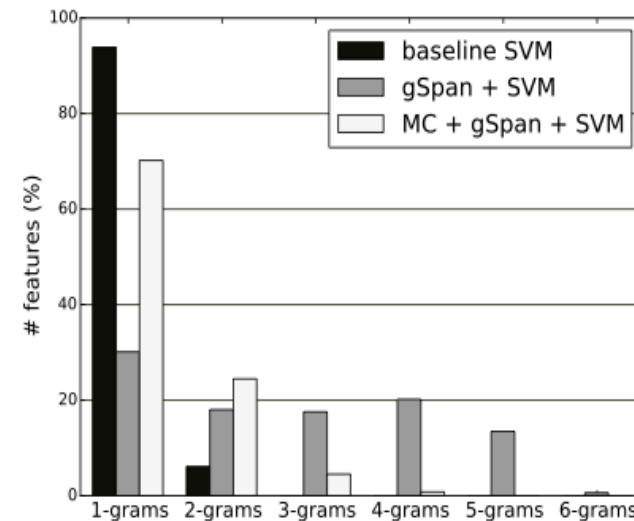


**Figure:** Number of subgraph features per support (%) on WebKB (left) and R8 (right) datasets. In black, the selected support chosen via the elbow method.

# Distribution of mined n-grams



**Figure:** Distribution of n-grams (standard and long-distance ones) among all the features on WebKB dataset.



**Figure:** Distribution of n-grams (standard and long-distance ones) among the top 5% most discriminative features for SVM on WebKB dataset.

# Contribution

- we explored a graph-of-words, to challenge the traditional bag-of-words for text classification.
- first trained a classifier using frequent sub-graphs as features for increased effectiveness.
- reduced each graph-of-words to its main core before mining the features for increased efficiency.
- reduced the total number of n-gram features considered in the baselines for little to no loss in prediction performances.

# GoW for text mining

- 1. Ad Hoc Information Retrieval
- 2. Single-Document Keyword Extraction
  - Elect *representative* words for a document that best describe it.
  - Sub-event Detection in Textual Streams (twitter)
  - keyword extraction Chrome extension
- 3. Text Categorization
  - as a Graph Classification Problem
  - Sentence based Kernels for short text categorization

# Sentence based Kernels for short text categorization [EMNLP 2015]

*Short text categorization*

Assume the sentences:

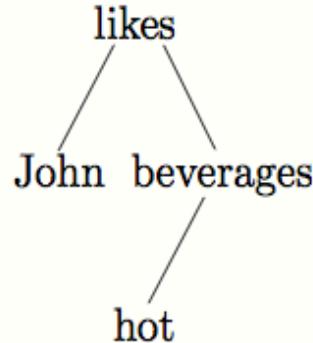
- ‘John likes hot beverages’
- ‘John loves warm drinks’

Their semantic similarity is obvious ...

Challenge: devise a similarity function that finds adequate similarity in these cases

# Dependency trees

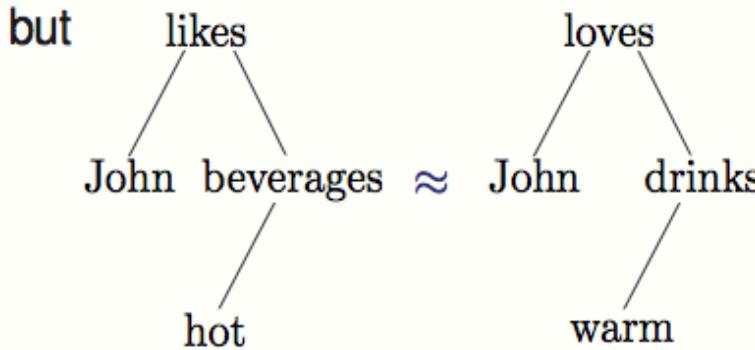
- Use of dependency tree structure to capture better bigrams



$$K_b('John likes hot beverages', 'John likes warm beverages') = 3 + 2$$

- Still problematic when inner nodes are different!

$$K_b('John likes hot beverages', 'John loves warm drinks') = 1$$



⇒ Sparsity issue for classification of *short documents* or *when little training data is available*

# Context

- Word embeddings
  - Embed words in a Low-dimension Euclidean vector space, suited for word similarity
  - Similar words are projects to near by places
  - word2vec: 300 dimensional representation of 3 million English words trained over a Google News dataset of 100 billion words using the Skip-gram model and a context size of 5
- Google word2vec [1]: query 'France', *distance* will display the most similar words and their distances to 'france'

Word	Cosine distance
spain	0.678515
belgium	0.665923
netherlands	0.652428
italy	0.633130
switzerland	0.622323
luxembourg	0.610033
portugal	0.577154
russia	0.571507
germany	0.563291
catalonia	0.534176

[1] <https://code.google.com/p/word2vec/>

# Mathematical framework (i)

- **Word Kernel (WK)**

We define a kernel between two words as a polynomial kernel - cosine similarity in the word embedding space, where  $\alpha$  is a scaling factor:

- **Phrase Kernel (PhK)**

$$WK(w_1, w_2) = \left[ \frac{1}{2} \left( 1 + \frac{\langle w_1, w_2 \rangle}{\|w_1\| \|w_2\|} \right) \right]^\alpha$$

We considered two types of phrases:

- *co-occurrence phrases*: contiguous sequences of words in the original sentence
- *syntactic phrases*: downward paths in the dependency tree representation

- **Product Kernel (PK)**

$w_{ij}$   $i$ -th word in phrase  $p_j$  of length  $i$

$$PK(p_1, p_2) = \prod_{i=1}^I WK(w_i^1, w_i^2)$$

- **Composition Kernel (CK)**

$$CK(p_1, p_2) = WK(\mathbf{p}_1, \mathbf{p}_2)$$

(in case of different sentence length) embedding  $\mathbf{p}_j$  of the phrase  $p_j$  obtained either by

- addition:

$$\mathbf{p}_j = (\sum_{i=1}^I \mathbf{w}_i^j)$$

- or element wise multiplication

$$\mathbf{p}_j = (\odot_{i=1}^I \mathbf{w}_i^j) \text{ i.e.}$$

i.e. ‘Berlin’ with ‘capital of Germany’

# Mathematical framework II

## ■ Sentence Kernel (SK)

- Defined through convolution as the sum of all local phrasal similarities, i. e. kernel values between phrases contained in the sentences:

$$\text{SK}(s_1, s_2) = \sum_{\substack{p_1 \in \phi(s_1), \\ p_2 \in \phi(s_2)}} \lambda_1^\epsilon \lambda_2^\eta \text{PhK}(p_1, p_2)$$

where

- $\phi(s_k)$ : set of phrases in sentence  $s_k$  (co-occurrence or syntactic phrases)
- $\lambda_1$ : decaying factor penalizing longer phrases
- $\epsilon = \max\{|p_1|, |p_2|\}$ : the maximum length of the two phrases
- $\lambda_2$ : distortion parameter
- $\eta = ||p_1| - |p_2||$ : length difference between the two phrases
- $\text{PhK}$ : phrase kernel (either PK, CK<sup>+</sup> or CK<sup>⊖</sup>)

## ■ Document Kernel

- Sum of all sentence kernel values between two documents
- We may normalize each kernel value  $K_{ij}$  between documents  $i$  and  $j$  by  $\sqrt{K_{ii}K_{jj}}$

# Experiments (i)

## ▪ Datasets

- *Binary sentiment analysis:*
  - Movie review dataset (PL05) of 10,662 sentences
  - Product review dataset (Amazon) of 2,000 multi-line documents for 4 different product groups
- *Ternary sentiment analysis:*
  - SemEval 2013 Task B dataset (Twitter) with 12,348 tweets
- *Binary subjectivity detection*
  - Movie review/summary dataset (PL04) of 10,000 sentences
  - Opinion corpus (MPQA) of 11,640 sentences
- *Seven-class topic spotting*
  - News dataset (News) of 32,602 one-line news summaries

## ▪ Experimental settings

- *FANSE parser* for dependency trees
- *LibSVM* for SVM training with *one-vs-one strategy* for multi-class tasks
- To prevent overfitting, we tuned the parameters using cross-validation on 80% of PL05 dataset and used the same set of parameters on the remaining datasets:
  - PK:  $\alpha = 5$ ,  $\lambda_1 = 1$  (no need for distortion as the phrases are of same length by definition)
  - CK:  $\alpha = 5$ ,  $\lambda_1 = \lambda_2 = 0.5$
- We performed normalization for our kernel and baselines *only when it led to performance improvements on the training set* (PL05, News, PL04 and MPQA)

# Experiments (ii)

## ■ Evaluation

- We report *accuracy* on:
  - PL05: remaining 20% test set
  - Twitter: standard test split (25%)
  - News: standard test split (50%)
  - Amazon, PL04, MPQA: 5-fold CV
- Significance test: *micro sign test* ( $p < 0.01$ ) over the bigram baseline with same phrase definition

## ■ Models

- Delta word kernel
  - standard unigram and bigram baseline approaches
- Dependency Tree Kernel (DTK)
  - $n$ -grams from downward paths in dependency trees
- Vector Tree Kernel (VTK)
  - $\phi(\cdot)$  enumerates all random walks in the dependency tree
- Our convolutional sentence kernels based on phrase kernels and polynomial word kernels

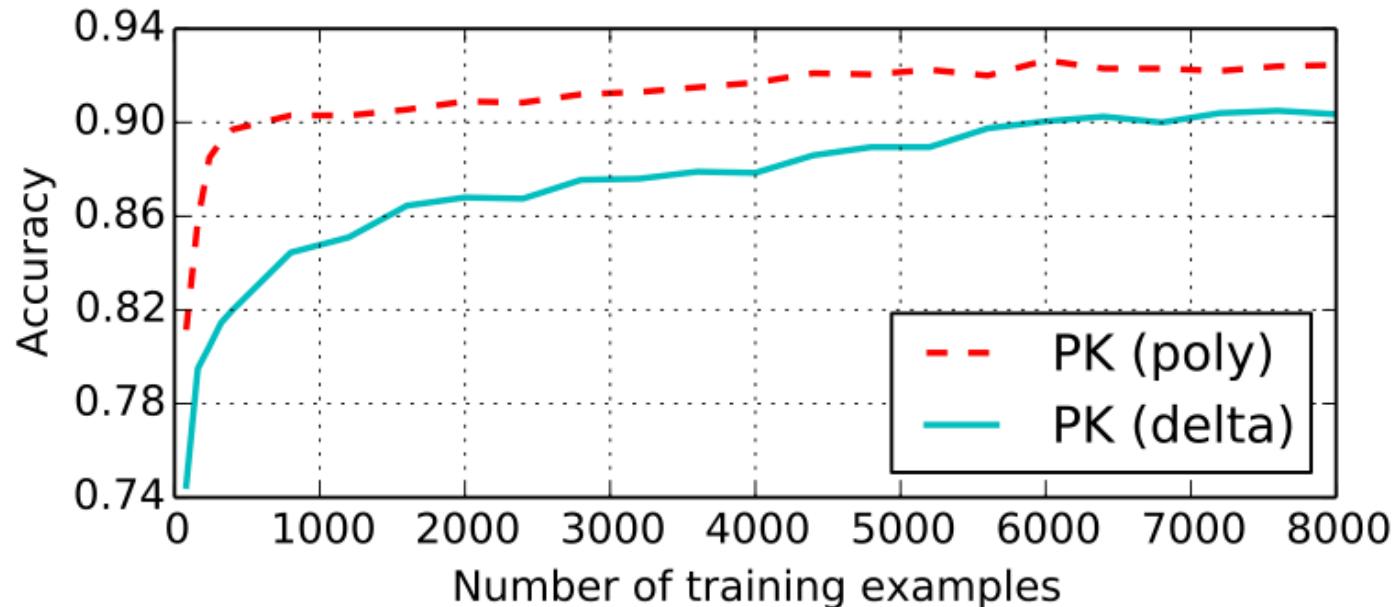
# Results (i)

- Resolving the sparsity issue for short documents

phrase definition	phrase kernel	phrase length	word kernel	PL05	Amazon	Twitter	News	PL04	MPQA
co-occurrence	PK	1	delta	0.742	0.768	0.623	0.769	0.904	0.754
co-occurrence	PK	2	delta	0.739	0.765	0.611	0.766	0.907	0.754
syntactic	PK	2	delta	0.748	0.791	0.646	0.767	0.910	0.757
random walk	PK	2	poly	0.799	0.810	0.698	0.802	<b>0.927</b>	<b>0.797</b>
co-occurrence	PK	1	poly	0.789*	0.797	0.776*	<b>0.806*</b>	0.923*	0.793*
co-occurrence	PK	2	poly	0.784*	0.798	0.762*	0.801*	0.926*	0.794*
co-occurrence	CK <sup>+</sup>	2	poly	0.796*	0.778	0.613	0.792*	0.917*	0.796*
co-occurrence	CK <sup>⊖</sup>	2	poly	<b>0.801*</b>	0.783	0.757*	0.793*	0.918*	0.794*
syntactic	PK	2	poly	0.796*	<b>0.813*</b>	<b>0.808*</b>	0.805*	<b>0.927*</b>	0.796*
syntactic	CK <sup>+</sup>	2	poly	0.794*	0.780	0.741*	0.788*	0.918*	0.794*
syntactic	CK <sup>⊖</sup>	2	poly	0.797*	0.774	0.744*	0.792*	0.918*	0.794*

# Results (ii)

- **Resolving the sparsity issue when little training data is available**
  - Learning done on 1% to 100% of the training set (80% of the dataset) on PL04
    - ⇒ our kernel starts to plateau earlier
    - ⇒ our kernel reaches the maximum baseline accuracy with only 1,500 training examples



# Related efforts

## Kernels

- *semantic kernels* for text categorization, capitalizing on WordNet: continuous word kernels based on the inverse of the path lengths in the tree rather than the common delta word kernel used so far, i.e. exact matching between unigrams [Siolas, d'Alché Buc. 2000]
- extensions to tree similarity measures [Bloehdorn et al. 2006]
- GVSM kernels [Mavroeidis, Vazirgiannis et.al 2005]
- *tree kernels* to compare trees based on their topology (e.g., shared subtrees) rather than the similarity between their nodes [Collins, Duffy 2001]
- Dependency Tree Kernels (DTK) to capture syntactic similarities [Culotta and Sorensen 2004]
- additional semantic meta-features based on word embeddings to enhance classification for short text similarity [Song and Roth (2015)[Kenter, Rijke (2015)]

## Word Embeddings

- Wordnet [Miller1995],
- SENNA[Collobert et al., 2011]
- Word2vec [Mikolov et al.2013]

# Conclusions

- sparsity issue classifying short documents or when little training data is available
- we proposed a novel convolutional sentence kernel based on word embeddings that tackles sparsity issue.
- achieved significant improvements over the baselines across all datasets when taking into account the additional information from the latent word similarity (word embeddings) and the syntactic structure (dependency tree).

# Graph of Words - Contributions

- Extraction – GoW for term weighting for IR, keyword extraction and classification
- Learning @ document level
  - new features - subgraphs of word for classification
- Regularization @ collection level
  - GoW - regularization as feature similarity (similar features should have similar weights)

## Future directions

- GoW as features to ML algorithms
- GoW for Document and Collection visualization
- Move from tw-idf to tw-**icw** and tw-**idw** ...
- Future work
- designing new kernels for syntactic parse trees – similarity between non-terminal nodes

# Novel models needed – Thought vectors

- “thought vectors”: Google is working on a new type of algorithm designed to encode thoughts as sequences of numbers
- Geoff Hinton – Google: “If you take the vector for Paris and subtract the vector for France and add Italy, you get Rome,” he said. “It’s quite remarkable.” @Guardian

<https://wtvox.com/robotics/google-is-working-on-a-new-algorithm-thought-vectors>

# Acknowledgements

## Collaborators

F. Rousseau(LIX), P. Meladianos (AUEB), Y. Nikoletzos (AUEB), E. Kiagias (AUEB), J. Kim (Oxford)

## Funding

Google, Airbus, DIGITEO, AXA (Chair)

# Relevant publications

- ① Rousseau, F. and M. Vazirgiannis (2013a). Composition of TF Normalizations: New Insights on Scoring Functions for Ad Hoc IR. In *Proceedings of ACM SIGIR '13*, pp. 917–920.
- ② Rousseau, F. and M. Vazirgiannis (2013b). Graph-of-word and TW-IDF: New Approach to Ad Hoc IR. In *Proceedings of the 22<sup>nd</sup> ACM CIKM '13*, pp. 59–68, *best paper mention award*.
- ③ Rousseau, F. and M. Vazirgiannis (2015a). Main Core Retention on Graph-of-words for Single-Document Keyword Extraction. In *Proceedings of the 37th European Conference on Information Retrieval. ECIR '15*.
- ④ P. Meladianos, Y. Nikolentzos, F. Rousseau, Y. Stavrakas and M. Vazirgiannis (2015b) Degeneracy-based Real-Time Sub-Event Detection in Twitter Stream. *Proceedings of the 9th AAAI International Conference on Web and Social Media (AAAI ICWSM '15)*.
- ⑤ Rousseau, F., Kiagias E. and M. Vazirgiannis, (2015c) Text Categorization as a Graph Classification Problem, in the proceedings of the ACL 2015 conference
- ⑥ Jonghoon Kim F. Rousseau M. Vazirgiannis, “Convolutional Sentence Kernel from Word Embeddings for Short Text Categorization”, EMNLP 2015 conference

# THANK YOU!

## Questions?

<http://www.lix.polytechnique.fr/dascim/>

**We are hiring!!**

ML for graph and text mining  
postdoc and PhD

Contact me: [mvazirg@lix.polytechnique.fr](mailto:mvazirg@lix.polytechnique.fr)